

## ACIT 3855 – Lab 9 – Issue Fixing

<b>Instructor</b>	Mike Mulder ( <a href="mailto:mmulder10@bcit.ca">mmulder10@bcit.ca</a> ) – Sets A and B Rafi Mohammad ( <a href="mailto:rafi_mohammad@bcit.ca">rafi_mohammad@bcit.ca</a> ) – Set C
<b>Total Marks</b>	10
<b>Due Dates</b>	Demo due by the end of next class: <ul style="list-style-type: none"><li>• Nov. 14<sup>th</sup> for Set A</li><li>• Nov. 16<sup>th</sup> for Sets B and C</li></ul>

### Purpose

- Fix the software issues we've encountered so far in building our software system
- Start to think about the process for troubleshooting and correcting software issues

**Make sure to test each of your fixes before moving on to the next.**

### Part 1 – Dashboard UI: Audit Log Indices Out of Synch with Data

Observed: The displayed Audit Log index and the corresponding Event are out of synch. The index is ahead of the data (i.e., the event data is for the previously displayed index).

Expected: The Audit Log index and data should be synchronized.

Solution: In the EndpointAudit.js file:

Add a state for the index:

```
const [index, setIndex] = useState(null);
```

Set the index upon successfully response from the audit endpoints:

```
setIndex(rand_val);
```

Update the view to display the index:

```
return (  
  <div>  
    <h3>{props.endpoint}-{index}</h3>  
    {JSON.stringify(log)}  
  </div>  
)
```

### Part 2 – Kafka Topics Aren't Persisting

This was intentionally done to allow for debugging and troubleshooting without having to deal with old data in the Kafka message queue.

Observed: The events topic in the Kafka broker does not persist when Kafka is brought back up, such as when you stop your VM, start it again and then run “docker-compose up -d”.

Expected: Kafka topics are meant to persist.

Solution: The kafka logs (i.e., where the messages on the topics are stored) are being stored on the container but are lost when the container is removed. The solution is to mount a volume or bind a VM folder to the container for the /kafka directory. In addition, the Zookeeper data also needs to persist to keep it in synch with Kafka.

Create the following two folders on your VM:

- /home/<username>/zookeeper/data
- /home/<username>/kafka

Add the following to your docker-compose.yml file.

For the zookeeper service add a volume (bind mount):

```
volumes:
  - /home/<username>/zookeeper/data:/opt/zookeeper-3.4.13/data
```

For the kafka service (bind mount):

```
volumes:
  - /var/run/docker.sock:/var/run/docker.sock
  - /home/<username>/kafka:/kafka/kafka-logs
```

Also add the KAFKA\_LOG\_DIRS and KAFKA\_BROKER\_ID to the environment for the kafka service:

```
environment:
  KAFKA_CREATE_TOPICS: "events:1:1" # topic:partitions
  KAFKA_ADVERTISED_HOST_NAME: system-acit3855.wes
  KAFKA_LISTENERS: INSIDE://:29092,OUTSIDE://:9092
  KAFKA_INTER_BROKER_LISTENER_NAME: INSIDE
  KAFKA_ADVERTISED_LISTENERS: INSIDE://kafka:29092
  KAFKA_LISTENER_SECURITY_PROTOCOL_MAP: INSIDE:PLAINTEXT
  KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181
  KAFKA_LOG_DIRS: /kafka/kafka-logs
  KAFKA_BROKER_ID: 1
```

This assigns a broker id to kafka (which it uses for the stored data) and the folder in which to store the logs data (which refers to the data in the kafka topics and partitions, not application logs).

Now when you docker-compose down and then docker-compose up -d your services, the messages on the events topic in Kafka will persist.

### Part 3 – Processing Service Stats are Inaccurate

Observed: The statistics for the number of events processed by your Processing Service are inaccurate due to the design of the Storage Service endpoints, which only take a single timestamp. So it returns all the events received since that timestamp, but the Processing Service doesn't know exactly the end time for those events and makes a "guess" by using its current time. That can result in a gap or overlap in the time range for which the events are retrieved.

Expected: The Processing Service should process exactly the number of events received by your system.

Solution: Add both a start and an end time to the Storage Service endpoints that retrieve events. So they will return all events created from a start time ( $\geq$  start\_time) and before an end time ( $<$  end\_time). This will make the statistics more accurate if the Processing Service queries for new events from its stored last updated timestamp and the current timestamp. Once it's retrieved the events, the Processing Service should store the current timestamp as the last updated timestamp.

In the Storage Service, update your two GET endpoints. Here is a sample of the code:

```
def get_blood_pressure_readings(start_timestamp, end_timestamp):
    """ Gets new blood pressure readings after the timestamp """

    session = DB_SESSION()

    start_timestamp_datetime =
        datetime.datetime.strptime(start_timestamp, "%Y-%m-%dT%H:%M:%S")
    end_timestamp_datetime =
        datetime.datetime.strptime(end_timestamp, "%Y-%m-%dT%H:%M:%S")

    results = session.query(BloodPressure).filter(
        and_(BloodPressure.date_created >= start_timestamp_datetime,
            BloodPressure.date_created < end_timestamp_datetime))
    ...
```

You'll also need an additional import: `from sqlalchemy import and_`. Make sure you also update the `openapi.yml` file accordingly with the new parameter.

In your Processing Service, use the last updated time and the current time to get your events. For example:

```
response = requests.get(app_config["eventstore"]["url"] +
                        "/blood-pressure?start_timestamp=" +
                        last_updated + "&end_timestamp=" +
                        current_timestamp)
```

Make sure to set the new last updated time in your SQLite database to the `current_timestamp` value so there are no gaps in your processing.

## Part 4 – Storage Service Does Not Connect on Startup

Observed: Storage Service sometimes does not connect to Kafka on startup, requiring the service to be stopped and restarted.

Expected: The Storage Service should be able to connect to Kafka on startup, even if it has to wait until Kafka is up and running.

Solution: The following code from the `process_messages` function is trying to connect to Kafka but is failing silently (no exception handling) and there is no re-try logic:

```
client = KafkaClient(hosts=hostname)  
topic = client.topics[str.encode(app_config["events"]["topic"])]
```

You should add re-try logic to this function. Here is some pseudo code:

Define a maximum number of retries (this should be obtained from your configuration file for the service)

Set the current retry count to zero.

While the current retry count is less than the maximum number of retries:

- Display an info log message indicating you are trying to connect to Kafka and the current retry count
- Create a `KafkaClient` and get the topic (i.e., the code above)
- If an exception is raised (hint: add exception handling - try/except - in the while loop):
  - Display an error log message indicating that the connection failed
  - Sleep for a configurable number of seconds (hint: import `time` and use `time.sleep`, obtain the sleep time for your configuration file)
  - Increment the current retry count

## Part 5 – (Improvement) Receiver Service Reconnects to Kafka for Each Event

This one you need to figure out yourself. Here is the approach you could take:

- Create the `KafkaClient` when your Receiver Service starts up, not in the functions for each endpoint.
- Add the re-try logic as in Part 4.
- The function for each endpoint will now just get a producer (i.e., `get_synch_producer`) and produce the event on the topic.

## Part 6 – (Bonus to First Student to Figure it Out) Storage Service Losing Connection

It appears that the Storage Service still may lose its Kafka connection after some time. We currently do not have a fix for this issue.

Observed: After a period of time, when `jMeter` is used to send in events to the system, the Storage Service will receive the first event and then not process any subsequent events until it is stopped and restarted.

Expected: The Storage Service should process events indefinitely until it is stopped or the Kafka broker goes down.

The first student who can come up with a solution and:

- Show the solution working on their system
- Mike Mulder can implement the solution in his system and determine that it solves the problem

Will get a small bonus on Assignment 2.

### Grading and Submission

Submit the following to the Lab 9 Dropbox on D2L:

- The app.py file for your Storage Service with the retry logic in process\_messages.
- The app.py file for your Receiver Service with the single Kafka connection.

Show the code for each of your solutions above.	5 marks
Demonstrate your system with jMeter proving the following: <ul style="list-style-type: none"><li>• The Dashboard UI shows the correct Audit Event data</li><li>• The events persist in the Kafka events topic after a docker-compose down followed by a docker-compose up -d</li><li>• The Processing Service has the correct statistics (within 5% of the expected)</li><li>• The Storage Service doesn't need to be restarted after a docker-compose down followed by a docker-compose up -d</li><li>• Receiver Service still works with one-time Kafka connection</li></ul>	5 marks
Storage Service connection loss	Bonus
<b>Total</b>	<b>10 marks</b>